

# Modular Bellman Dynamic Programming

## Dyn-X: Theory to High-Dimensional Implementation

Chris Carroll (*JHU*)    Alan Lujan (*JHU*)

Akshay Shanker (*UNSW/Econ-ARK*)    Matt White (*Econ-ARK*)

CEF 2025

July 2025

# Roadmap

---

1. **Problem:** Monolithic DP code
2. **Core Idea:** Factored Bellman  $\rightarrow$  Modular CDA decomposition
3. **Implementation:** `perches`, `movers`, and `stages`
4. **Case Study:** Housing model with branching
5. **Future:** Implementation, open questions & roadmap

We have clean and \*\*modular theory\*\* with all primitives.


$$v_{t+1} = \mathbb{T}v_t, \quad v_t, v_t \in \mathcal{B}(X)$$

What happens when we are tasked to compute and estimate models with policy or business applications specific features?

# Why Is Dynamic-Programming Code Still Monolithic?

## Introduction

- **Ad hoc, one-off, copy-paste implementations** → near-zero reuse.
- **Curse of dimensionality** × **curse of idiosyncratic engineering**: each new model is bespoke with handcrafted hacks.
- **Implications for supercomputing**: Solution structure matters.
- *Powell (2012)*: introducing a post-decision state can "*dramatically simplify*" dynamic programs — yet frameworks rarely expose it.



```
# All-in-one nested loops
for t in range(T-1, -1, -1):
    for i, a in enumerate(asset_grid):
        best_val = -np.inf
        for c in np.linspace(0, a, 101):
            # Utility + transition mixed
            u = np.sqrt(c)
            a_next = a - c
            j = int(round(a_next))
            # Bellman buried in loops
            val = u + beta * V[t+1, j]
            if val > best_val:
                best_val = val
        V[t, i] = best_val
```

### Extreme Monolith:

- Triple nested loops
- State dynamics entangled with optimization
- Hard to modify or extend
- No real operator representation

## Our premise

*Operator theoretic approach **reduces mathematical complexity**, it should also reduce computational complexity.*

```
# Separate operators for each decision
def cntn_to_dcsn(EV_next, model):
    # Backward: optimize consumption
    ...
    return V_dcsn

def dcsn_to_arvl(V_dcsn, model):
    # Backward: integrate over shocks
    ...
    return V_arvl

# Clean operator composition
T_stage = dcsn_to_arvl(cntn_to_dcsn(...))
```

Bellman factored:  $\mathcal{T}^F = \mathcal{T}^{va} \circ \mathcal{T}^{ev}$

### Main ideas:

- Recursive problem decomposed into three operations.
- Explicitly define operations as flexible computational objects and **remove reference to time** within stage.
- Compose operators into canonical stages.
- Dynamic program turns out to be a **graph** where operator arguments are nodes (perches) and operations are edges (movers)

**Primal:** Mixes action with future shock

$$V_{t+1}(x) = \max_{a \in A(x)} \left\{ r(x, a) + \beta \mathbb{E} [V_t(f(x, a, W_{t+1}))] \right\}$$

where  $V_{t+1}$  is the value function at time  $t + 1$  and  $V_t$  is the value function at time  $t$ ,  $A(x)$  is the action set at state  $x$ ,  $r(x, a)$  is the reward function,  $f(x, a, W_{t+1})$  is the transition function, and  $W_{t+1}$  is the shock at time  $t + 1$ .

**Factored:** Separate transitions and continuation value

$$x_e = g_{ve}(x_v, a), \quad x'_v = g_{av}(x_e, W'), \quad \mathcal{E}(x_e) = \mathbb{E}_{W'}[V(x'_v)]$$

$$V(x_v) = \max_a \{ r(x_v, a) + \beta \mathcal{E}(g_{ve}(x_v, a)) \}$$

Key: Post-Decision State

$x_e$  = Powell's post-decision state — no nested expectation inside max



Label	Timing	Carries value
Arrival $x_a$	start (pre-shock)	$\mathcal{A}(x_a)$
Decision $x_v$	post-shock	$\mathcal{V}(x_v)$
Continuation $x_e$	post-action	$\mathcal{E}(x_e)$

Operators:  $\mathcal{T}^{ev}$  (optimise),  $\mathcal{T}^{va}$  (expectation).

### Mathematical Decomposition:

$$\mathcal{T}^F = \mathcal{T}^{va} \circ \mathcal{T}^{ev}$$

$$\begin{cases} (\mathcal{T}^{ev} \mathcal{E})(x_v) = \max_a \{r + \beta \mathcal{E}\} \\ (\mathcal{T}^{va} \mathcal{V})(x_a) = \mathbb{E}_W[\mathcal{V}(g_{av})] \end{cases}$$

### Computational Needs:

- Store functions & policies
- Swap solvers (VFI/EGM)
- Connect stages
- Enable simulation

## Solution

Map mathematical objects directly to computational structures

## Perches (nodes):

- arvl:  $\mathcal{A}, \mu_a$
- dcsn:  $\mathcal{V}$ , policy,  $\mu_v$
- cntn:  $\mathcal{E}, \mu_e$

## Movers (edges):


- Backward:  $\mathcal{T}^{ev}, \mathcal{T}^{va}$
- Forward: push distributions
- Swap solvers via config

Each perch stores .sol, .dist, .grid

A stage is a union of two direct acyclic graphs (DAGs) with `perches` as nodes and `movers` as edges.

```
arvl --shock--> dcsn --action--> cntn
^                               |
|----- Bellman backward -----|
```

Graph is the composition  $\mathcal{T}^{va} \circ \mathcal{T}^{ev}$ .



Dashed = Bellman (value) recursion; solid = push-forward of measures. This template underlies every Dyn-X stage.

1. Load YAML  $\Rightarrow$  symbolic model
2. Compile  $\Rightarrow$  NumPy/JIT functions & grids
3. Solve backward via **Horse** (solver engine) + **Whisperer** (orchestrator)
4. Simulate forward (horse)

## Clean Separation

- **Model specification** (YAML/config) is separate from **solution algorithm**
- Horse = pluggable solver (VFI, EGM, etc.)
- Whisperer = coordinates solving across stages

- **Thin operators** each mover  $\approx$  one NumPy/Numba kernel  $\Rightarrow$  launch exactly the same kernel on a GPU array
- **No global state** per-stage data live in `perch.sol`  $\Rightarrow$  host  $\leftrightarrow$  device copy is local, automated
- **Zero code change** Switch backends via config: CPU  $\leftrightarrow$  GPU  $\leftrightarrow$  MPI  
Unlike ad-hoc GPU implementations
- **Tailored kernels** Custom CUDA for Bellman max (not generic autodiff)  
Handles non-differentiable policies, discrete choices

Outcome in the housing benchmark (10k wealth  $\times$  30 housing  $\times$  5 income points):

run-time: 74 s (CPU)  $\rightarrow$  2.6 s (1 V100)

**$\sim 30\times$  speedup** — comparable to specialized GPU codes, but modular

```
1 state_space:
2   decision: [assets, y, house]
3   arrival:  [assets, house]
4 functions:
5   utility: "alpha*np.log(c)+(1-alpha)*np.log(kappa*(house+1))"
```

StageCraft infers the CDA graph automatically.



1. TENU — tenure choice (branching)
2. OWNH — owner housing stock
3. RNTH — renter housing services
4. OWNC — owner consumption (EGM)
5. RNTC — renter consumption (EGM)

# Branching Bellman — Tenure Choice

## Case Study: Housing Model

---

Decision state  $(a, H, y)$ :

$$V_v(a, H, y) = \max \left\{ V_e^{\text{rent}}((1+r)a + H, y), V_e^{\text{own}}(a, y, H) \right\}.$$

Arrival update:

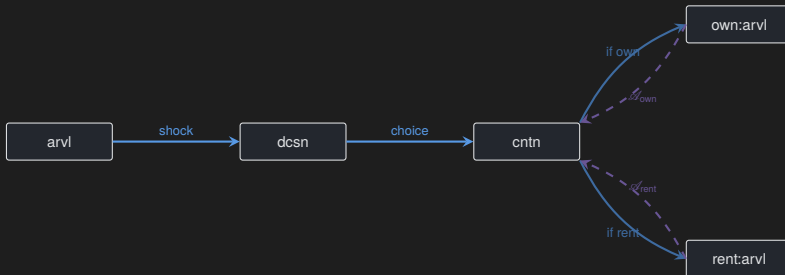
$$V_a(a, H, \bar{y}) = \mathbb{E}_{\xi} [V_v(a, H, f(\bar{y}, \xi))].$$

- Rent branch liquidates housing  $\Rightarrow w = (1+r)a + H$ .
- Both continuations stored in `cntn.sol` under keys "rent", "own".

# Branching Stage Graph (single cntn perch)

## Housing Model

Case Study:



Forward: Single path within stage, branches after `cntn`.

Backward: Multiple value functions aggregate at `cntn`.

Decision state  $(w, y)$  with discrete grid  $\mathbb{H}$ :

$$\mathcal{T}^{ev} \mathcal{E}(w, y) = \max_{\substack{S \in \mathbb{H} \\ P^r S \leq w}} \mathcal{E}(S, y, w - P^r S).$$

Budget feasibility  $P^r S \leq w$  is enforced in the choice set.

## YAML Snippet — Discrete Choice Case Study: Housing Model

```
1 dcsn_to_cntn:
2   type: forward
3   method: discrete_choice
4   choice_grid:
5     S: {type: linspace, min: S_min, max: S_max, points: S_pts}
6   constraints:
7     - "P_r*S <= w"
```

The backward `cntn_to_dcsn` mover enumerates the  $S$ -grid (stored in `cntn.grid["S"]`) and takes the element-wise max.

### Solver Methods:

- **VFI** + GPU (30× speedup)
- **EGM** (Carroll 2006)
- **DCEGM** (Iskhakov et al. 2017)
- **FUES** (4× fewer grid points)

### Forward Simulation:

`arvl.dist` → `dcns.dist` →  
`cntn.dist`

Mass preserved at each step

Key: Swap solvers via config, not code

### Generic ML Frameworks:

- TensorFlow/JAX: Arrays & ops
- No economic primitives
- Generic autodiff  $\neq$  EGM

### Economics Libraries:

- HARK: Modular but not graph
- Dolo: DSL but no factoring
- QuantEcon: Solvers not stages

**Dyn-X: Economic objects as nodes, operators as edges**

### Core conjectures:

- **Universality:** Any MDP  $\rightarrow$  CDA form?
- **Minimality:** Are irreducible stages truly minimal?
- **Edge cases:** Simultaneous shocks, non-period models

### Scaling up:

- **Heterogeneous agents:** Multiple circuits in parallel
- **General equilibrium:** Link circuits via market clearing
- **Infinite horizon:** Period structure with stationary shocks

Empirical validation: 30 $\times$  speedup on housing model



- **GPU via Numba CUDA** (implemented, 30× speedup)
- **MPI** for distributed computing
- **JAX/cuPy backends** (planned)
- **Lean-4 export** for formal verification

Dyn-X 1.0 beta: Q4 2025 — [github.com/econ-ark/dynx](https://github.com/econ-ark/dynx)

- **Problem:** Monolithic DP code — no reusability
- **Solution:** Factored Bellman → CDA stages
- **Implementation:** Graph of economic objects
- **Performance:** 30× speedup with GPU

## Impact

Focus on economics, not implementation

Dyn-X 1.0 beta: Q4 2025

Questions & discussion.

Slides & code → [github.com/akshay-shanker/dynx-cef-2025](https://github.com/akshay-shanker/dynx-cef-2025)

# Modular MDP — Formal Definition

## Tuple

$$((\Omega, \Sigma, \mathbb{P}), J, \{\mathcal{X}_j\}, \{\mathcal{F}_{\mathcal{X}_j}\}, \{\mathfrak{S}_j, \tilde{\mathfrak{S}}_j\}, \{\text{CDC}_j\})$$

1.  $J$ : stage indices (finite or countable)
2.  $\mathcal{X}_j$ : topological vector space of states
3.  $\mathcal{F}_{\mathcal{X}_j} \subseteq \mathcal{M}(\mathcal{X}_j, \mathbb{R})$
4. Shocks  $\mathfrak{S}_j$  (observable) and  $\tilde{\mathfrak{S}}_j$  (latent)
5.  $\text{CDC}_j : \mathcal{F}_{\mathcal{X}_{j+1}} \rightarrow \mathcal{F}_{\mathcal{X}_j}$

Stage  $j$  is *irreducible* if  $\sigma(\mathfrak{S}_j)$  and  $\sigma(\tilde{\mathfrak{S}}_j)$  admit no non-trivial independent sub-algebras.

# Complex Example: Mortgage Model with Sub-Circuits

